

# ML\_1

November 10, 2025

```
[102]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.ensemble import RandomForestRegressor
from sklearn.preprocessing import StandardScaler
from sklearn.metrics import r2_score, mean_squared_error, mean_absolute_error
```

```
[103]: # Load the dataset (change path if needed)
df = pd.read_csv(r"C:\Users\suraj\OneDrive\Desktop\LP3-master\ML\datasets\uber.csv")
# Display first 5 rows
df.head()
```

```
[103]:   Unnamed: 0                               key  fare_amount \
0    24238194    2015-05-07 19:52:06.0000003      7.5
1    27835199    2009-07-17 20:04:56.0000002      7.7
2    44984355    2009-08-24 21:45:00.00000061     12.9
3    25894730    2009-06-26 08:22:21.0000001      5.3
4    17610152    2014-08-28 17:47:00.000000188     16.0

                                         pickup_datetime  pickup_longitude  pickup_latitude \
0    2015-05-07 19:52:06 UTC                  -73.999817        40.738354
1    2009-07-17 20:04:56 UTC                  -73.994355        40.728225
2    2009-08-24 21:45:00 UTC                  -74.005043        40.740770
3    2009-06-26 08:22:21 UTC                  -73.976124        40.790844
4    2014-08-28 17:47:00 UTC                  -73.925023        40.744085

  dropoff_longitude  dropoff_latitude  passenger_count
0            -73.999512        40.723217             1
1            -73.994710        40.750325             1
2            -73.962565        40.772647             1
3            -73.965316        40.803349             3
4            -73.973082        40.761247             5
```

```
[104]: # Remove missing values
df = df.dropna()

# Keep only positive fare values
df = df[df['fare_amount'] > 0]

# Convert pickup_datetime to datetime format
df['pickup_datetime'] = pd.to_datetime(df['pickup_datetime'], errors='coerce')
df = df.dropna(subset=['pickup_datetime'])

# Create a simple distance feature (approximation)
df['distance'] = ((df['dropoff_longitude'] - df['pickup_longitude'])**2 +
                   (df['dropoff_latitude'] - df['pickup_latitude'])**2)**0.5

# Remove rows where distance is 0
df = df[df['distance'] > 0]

print("Data after preprocessing:", df.shape)
df.head()
```

Data after preprocessing: (194347, 10)

```
[104]: Unnamed: 0 key fare_amount \
0 24238194 2015-05-07 19:52:06.0000003 7.5
1 27835199 2009-07-17 20:04:56.0000002 7.7
2 44984355 2009-08-24 21:45:00.00000061 12.9
3 25894730 2009-06-26 08:22:21.0000001 5.3
4 17610152 2014-08-28 17:47:00.000000188 16.0

pickup_datetime pickup_longitude pickup_latitude \
0 2015-05-07 19:52:06+00:00 -73.999817 40.738354
1 2009-07-17 20:04:56+00:00 -73.994355 40.728225
2 2009-08-24 21:45:00+00:00 -74.005043 40.740770
3 2009-06-26 08:22:21+00:00 -73.976124 40.790844
4 2014-08-28 17:47:00+00:00 -73.925023 40.744085

dropoff_longitude dropoff_latitude passenger_count distance
0 -73.999512 40.723217 1 0.015140
1 -73.994710 40.750325 1 0.022103
2 -73.962565 40.772647 1 0.053109
3 -73.965316 40.803349 3 0.016528
4 -73.973082 40.761247 5 0.051031
```

```
[105]: # Remove extreme outliers using IQR method on fare_amount
Q1 = df['fare_amount'].quantile(0.25)
Q3 = df['fare_amount'].quantile(0.75)
IQR = Q3 - Q1
low = Q1 - 1.5 * IQR
```

```

high = Q3 + 1.5 * IQR
df = df[(df['fare_amount'] >= low) & (df['fare_amount'] <= high)]

# Remove unrealistic distance values (greater than 5 degrees ~ 500 km)
df = df[df['distance'] < 0.04]

print("Data after removing outliers:", df.shape)
df.head()

```

Data after removing outliers: (146319, 10)

```

[105]:      Unnamed: 0                 key  fare_amount \
0     24238194  2015-05-07 19:52:06.0000003      7.5
1     27835199  2009-07-17 20:04:56.0000002      7.7
3     25894730  2009-06-26 08:22:21.0000001      5.3
8     15822268  2012-02-17 09:32:00.00000043     9.7
10    2205147   2015-05-22 17:32:27.0000004      6.5

                  pickup_datetime  pickup_longitude  pickup_latitude \
0  2015-05-07 19:52:06+00:00          -73.999817        40.738354
1  2009-07-17 20:04:56+00:00          -73.994355        40.728225
3  2009-06-26 08:22:21+00:00          -73.976124        40.790844
8  2012-02-17 09:32:00+00:00          -73.975187        40.745767
10 2015-05-22 17:32:27+00:00          -73.974388        40.746952

  dropoff_longitude  dropoff_latitude  passenger_count  distance
0      -73.999512       40.723217             1  0.015140
1      -73.994710       40.750325             1  0.022103
3      -73.965316       40.803349             3  0.016528
8      -74.002720       40.743537             1  0.027623
10     -73.988586       40.729805            1  0.022262

```

```

[106]: # Select numerical columns for correlation
corr = df[['fare_amount', 'distance']].corr()

# Print correlation
print("Correlation Matrix:")
print(corr)

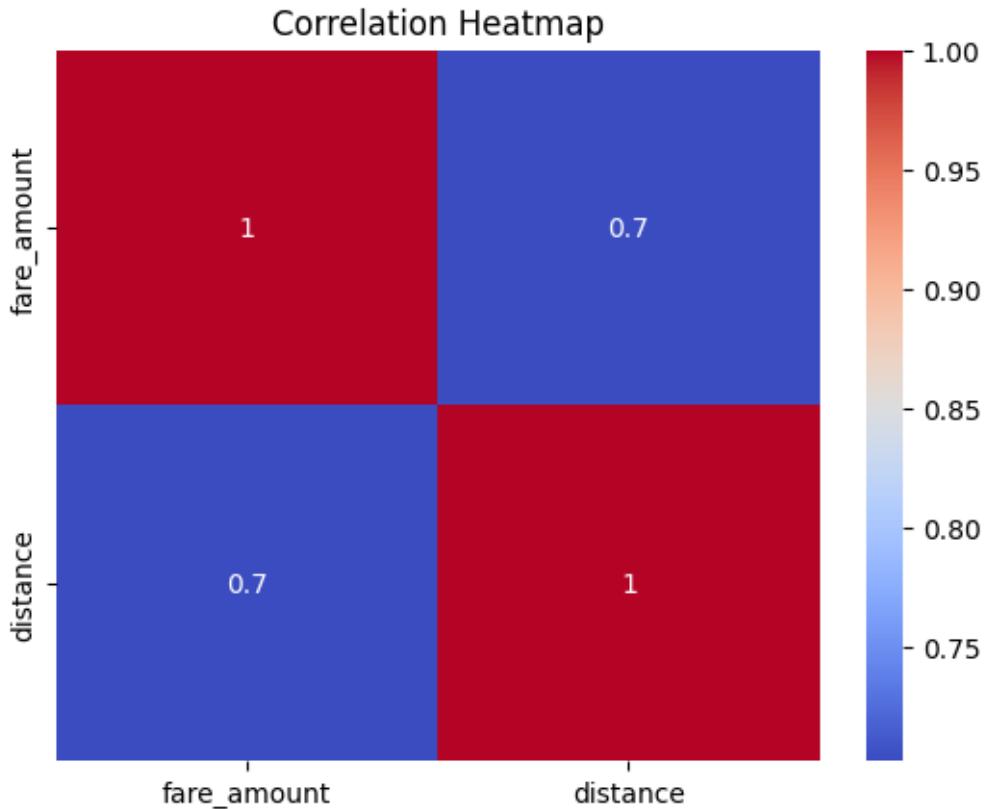
# Plot heatmap
sns.heatmap(corr, annot=True, cmap='coolwarm')
plt.title("Correlation Heatmap")
plt.show()

```

```

Correlation Matrix:
      fare_amount  distance
fare_amount      1.000000  0.702447
distance        0.702447  1.000000

```



```
[107]: # Extract time-based features
df['hour'] = df['pickup_datetime'].dt.hour
df['weekday'] = df['pickup_datetime'].dt.weekday
df['month'] = df['pickup_datetime'].dt.month

# Keep only useful columns
df = df[['fare_amount', 'distance', 'hour', 'weekday', 'month']]
df.head()
```

```
[107]:    fare_amount  distance  hour  weekday  month
0          7.5  0.015140    19        3      5
1          7.7  0.022103    20        4      7
3          5.3  0.016528     8        4      6
8          9.7  0.027623    9        4      2
10         6.5  0.022262   17        4      5
```

```
[108]: corr = df.corr()
print(corr)
```

	fare_amount	distance	hour	weekday	month
fare_amount	1.000000	0.702447	0.026125	-0.004165	0.038995

```

distance      0.702447  1.000000  0.013780  0.023288  0.007521
hour         0.026125  0.013780  1.000000 -0.076056 -0.002268
weekday     -0.004165  0.023288 -0.076056  1.000000 -0.010123
month        0.038995  0.007521 -0.002268 -0.010123  1.000000

```

```
[109]: # Define features and target
X = df[['distance', 'hour', 'weekday', 'month']]
y = df['fare_amount']

# Split dataset
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

print("Training data:", X_train.shape)
print("Testing data:", X_test.shape)
```

Training data: (117055, 4)  
 Testing data: (29264, 4)

```
[110]: # Scale features for Linear Regression
scaler = StandardScaler()
X_train_s = scaler.fit_transform(X_train)
X_test_s = scaler.transform(X_test)

# Train Linear Regression
lr = LinearRegression()
lr.fit(X_train_s, y_train)

# Predict
y_pred_lr = lr.predict(X_test_s)
```

```
[111]: # Train Random Forest Regressor
rf = RandomForestRegressor(n_estimators=100, random_state=42)
rf.fit(X_train_s, y_train)

# Predict
y_pred_rf = rf.predict(X_test_s)
```

```
[123]: # Define evaluation function
def evaluate(y_true, y_pred):
    r2 = r2_score(y_true, y_pred)
    rmse = np.sqrt(mean_squared_error(y_true, y_pred))
    mae = mean_absolute_error(y_true, y_pred)
    return round(r2, 4), round(rmse, 2), round(mae, 2)

# Evaluate both models
r2_lr, rmse_lr, mae_lr = evaluate(y_test, y_pred_lr)
r2_rf, rmse_rf, mae_rf = evaluate(y_test, y_pred_rf)
```

```
# Create comparison table
results = pd.DataFrame({
    'Model': ['Linear Regression', 'Random Forest'],
    'R2 Score': [r2_lr, r2_rf],
    'RMSE': [rmse_lr, rmse_rf],
    'MAE': [mae_lr, mae_rf]
})

print("Model Comparison Results:")
display(results)
```

Model Comparison Results:

	Model	R2 Score	RMSE	MAE
0	Linear Regression	0.4982	2.09	1.47
1	Random Forest	0.4700	2.15	1.53

[ ]: