

# ML\_4

November 11, 2025

```
[1]: import pandas as pd
import numpy as np
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score, confusion_matrix, precision_score, recall_score, classification_report
from sklearn.preprocessing import StandardScaler
from sklearn.neighbors import KNeighborsClassifier
```

```
[3]: df = pd.read_csv(r"C:\Users\suraj\OneDrive\Desktop\LP3-master\ML\datasets\diabetes.csv")
```

```
[5]: df.head()
```

```
[5]:    Pregnancies  Glucose  BloodPressure  SkinThickness  Insulin  BMI \
0            6        148            72            35         0  33.6
1            1         85            66            29         0  26.6
2            8        183            64             0         0  23.3
3            1         89            66            23         94  28.1
4            0        137            40            35        168  43.1
```

```
[5]:    Pedigree  Age  Outcome
0      0.627   50       1
1      0.351   31       0
2      0.672   32       1
3      0.167   21       0
4      2.288   33       1
```

```
[7]: df.shape
```

```
[7]: (768, 9)
```

```
[9]: df.isna().sum()
```

```
[9]: Pregnancies      0
Glucose          0
BloodPressure     0
SkinThickness     0
Insulin          0
```

```
BMI          0  
Pedigree    0  
Age          0  
Outcome     0  
dtype: int64
```

```
[11]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 768 entries, 0 to 767  
Data columns (total 9 columns):  
 #   Column        Non-Null Count  Dtype     
---  --  
 0   Pregnancies  768 non-null    int64    
 1   Glucose       768 non-null    int64    
 2   BloodPressure 768 non-null    int64    
 3   SkinThickness 768 non-null    int64    
 4   Insulin       768 non-null    int64    
 5   BMI           768 non-null    float64  
 6   Pedigree      768 non-null    float64  
 7   Age           768 non-null    int64    
 8   Outcome       768 non-null    int64    
dtypes: float64(2), int64(7)  
memory usage: 54.1 KB
```

```
[13]: X = df.drop('Outcome', axis=1)  
y = df['Outcome']
```

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.  
                                                 ↪2, random_state=42)
```

```
[15]: scaler = StandardScaler()  
  
X_train_scaled = scaler.fit_transform(X_train)  
X_test_scaled = scaler.transform(X_test)  
  
knn = KNeighborsClassifier(n_neighbors=5)  
knn.fit(X_train_scaled, y_train)  
  
y_pred = knn.predict(X_test_scaled)
```

```
[17]: cm = confusion_matrix(y_test, y_pred)  
print("confusion matrix:\n", cm)
```

```
confusion matrix:  
[[79 20]  
 [27 28]]
```

```
[24]: accuracy = accuracy_score(y_test, y_pred)
error_rate = 1-accuracy
precision = precision_score(y_test, y_pred)
recall = recall_score(y_test, y_pred)

print(f"Accuracy :{accuracy:.4f}")
print(f"Error rate : {error_rate: .4f} ")
print(f"Precision : {precision: .4f}")
print(f"Recall : {recall: .4f}")
```

Accuracy :0.6948  
Error rate : 0.3052  
Precision : 0.5833  
Recall : 0.5091

```
[26]: print("KNN Classification Report:\n", classification_report(y_test, y_pred, zero_division=0))
```

KNN Classification Report:

	precision	recall	f1-score	support
0	0.75	0.80	0.77	99
1	0.58	0.51	0.54	55
accuracy			0.69	154
macro avg	0.66	0.65	0.66	154
weighted avg	0.69	0.69	0.69	154

[ ]: